

# The Evolution of the "Turing Equation" Toward Perfection

## Recap: The Initial Turing Equation Framework

The journey began with a **symbolic formula** meant to capture an *auto-evolving learning process*. In its initial form, the so-called "**Equação de Turing**" was given as:

$$\sum_i (\alpha_i \times \beta_i) + \lambda \cdot f(\sum_t \phi_t)^{\infty}$$

This compact expression encodes a **meta-learning loop**:

- **$\alpha_i$  (Adaptive Capacity)**: measures the learner's slope of improvement (learning progress) on component  $i$  (e.g. change in performance per training steps, measured over a sliding window). It reflects *how much was learned* <sup>1</sup>.
- **$\beta_i$  (Task Complexity)**: a normalized difficulty of task/component  $i$  (ranging 0–1). It ensures learning is challenged by non-trivial problems.
- **$\lambda$  (Meta-Rate Constant)**: a small positive constant (e.g. 0.1) driving intrinsic motivation. This term encourages exploration and avoids stagnation by providing a baseline progress incentive even when  $\alpha$  or  $\beta$  are low.
- **$f(\sum_t \phi_t)^{\infty}$  (Recursive Resultant Function)**: a function  $f$  that takes the *accumulated changes*  $\sum_t \phi_t$  over time  $t$  and feeds it back into itself indefinitely (denoted by the exponent  $\infty$ ). In other words, the output influences future inputs, creating a **never-ending feedback loop** of learning. The system never fully converges; it continually evolves its state  $\phi_t$  (which could represent knowledge or parameters) over an infinite horizon.

In plainer terms, the **Turing Equation** tries to formalize an **AI that learns forever**: at each iteration it selects tasks with some difficulty  $\beta$ , makes progress  $\alpha$ , adds a constant drive  $\lambda$ , and uses a function  $f$  to update itself with the experience  $\phi_t$ , repeating this *ad infinitum*. This aligns with Alan Turing's intuition that machines could *"learn through trial and error, thereby improving their intelligence over time, much like humans"* <sup>2</sup>.

## The Student-Teacher Learning Loop

The equation above was operationalized via a **student-teacher framework**:

- **Student Model  $\mathcal{M}(\phi)$** : an AI model with parameters  $\phi$  that learns from tasks. This is the "Aluno" which improves over time.
- **Task Generator  $\mathcal{G}(\psi)$** : a "teacher" or curriculum policy with parameters  $\psi$  that proposes tasks (or environment variations) for the student. It decides *which task to present next*.
- **Replay Buffer  $\mathcal{R}$** : a memory of past tasks that yielded high learning progress (to be revisited later). This prevents forgetting and ensures valuable experiences are not lost.

The **outer meta-loop** runs as follows (per iteration):

1. **Sample a task  $\tau$  from  $\mathcal{G}(\psi)$** : The generator proposes a task with certain difficulty  $\beta(\tau)$ . Initially, tasks are sampled broadly since  $\mathcal{G}$  is not yet trained.
2. **Train  $\mathcal{M}(\phi)$  on  $\tau$** : The student model trains for  $n$  steps on this task.
3. **Measure performance change**: Compute the metric (e.g. accuracy or reward) *before* ( $p_{\text{before}}$ ) and *after* ( $p_{\text{after}}$ ) training on  $\tau$ . The **learning progress** on  $\tau$  is  $\alpha = p_{\text{after}} - p_{\text{before}}$ .
4. **Update score  $s$  for the task**: Using the Turing Equation's core rule, calculate  $s = \alpha \cdot \beta(\tau) + \lambda$ , where  $\alpha$  is the measured progress

and  $\beta(\tau)$  is the task's difficulty. This  $s$  serves as a **reward signal** for the task generator <sup>1</sup>. Intuitively, tasks that are *just hard enough* to yield progress score highest <sup>3</sup> (easy tasks give low  $\beta$ , overly hard tasks give low  $\alpha$ ; the sweet spot is the zone of proximal development where learning progress is maximized). 5. **Train the generator  $G(\psi)$** : The generator's parameters are updated (e.g. via policy gradient or bandit algorithms) to increase the probability of sampling tasks with higher expected  $s$ . *In effect, the teacher learns to present tasks that maximize the student's learning progress* <sup>1</sup>. 6. **Update Replay Buffer**: If task  $\tau$  yielded high progress (a positive and significant), add it to the replay buffer  $R$  (with priority weight related to  $s$ ). If a task yields negligible progress consistently, it may be retired from  $R$ . 7. **Inner loop mix (student training)**: In addition to new tasks, the student periodically trains on a mix of replayed tasks from  $R$  (e.g. 20% replay, 80% new tasks) to reinforce past learning and prevent forgetting. 8. **Repeat indefinitely**: Over many meta-iterations,  $G(\psi)$  adaptively focuses the curriculum,  $M(\phi)$  grows more capable, and the cycle continues without a fixed end – embodying the  $f(\dots)^{\infty}$  feedback in the equation.

This process **automates curriculum learning**, an idea supported by research showing that a teacher algorithm can sequentially select tasks to maximize a student's improvement <sup>1</sup> <sup>4</sup>. It aims for **open-ended learning**: the student and tasks co-evolve, constantly pushing the frontier of what the student can do.

## Challenges and Refinements in the Meta-Loop

As the team developed this system (both theoretically and in code), they identified several challenges or “bugs” that needed refinement:

- **Accurate Progress Baseline**: Initially, the learning progress  $\alpha$  was computed using the *latest performance* as part of the historical average, inadvertently **leaking information** and underestimating true progress. The fix was to compute the baseline performance **before** adding the new result. In practice, maintain a moving window of past performance  $p_{t-k}$  for each task/context and compute  $\alpha = p_t - \text{mean}(p_{t-k})$  with the new  $p_t$  excluded from the average. This gives a more stable, true improvement signal.
- **Stable Task Difficulty Mapping**: Originally, each task sample came with a freshly randomized underlying problem (e.g. a new random weight matrix in a synthetic task), meaning the student faced entirely new functions each time. This made it hard to measure progress, since no continuity existed. The solution was to **tie each difficulty level to a stable underlying task distribution** – for example, use a fixed random seed or a fixed target function per difficulty bin. This way, improvement can accumulate on tasks of each type, and  $\alpha$  reflects learning rather than the noise of a brand new task definition.
- **Device Consistency**: Small engineering fixes were needed (especially in code) to ensure that if the student model runs on GPU, the task data and generator policy also operate on the same device. Misaligned devices caused runtime errors but were straightforward to correct (e.g. by moving generated data tensors to GPU and registering difficulty parameters properly in  $G(\psi)$ ).
- **Replay Integration**: The initial loop trained the student on new tasks first and then did a separate mini-training on replay tasks, which made it hard to guarantee the intended 80/20 mix. The fix was to **construct a single combined batch** containing (for example) 80% new task data and 20% replay task data, and train the student on that in one go. This directly enforces the mix and ensures the student sees a blend of new challenges and important past ones.
- **Entropy and Exploration**: Without constraints, the task generator  $G(\psi)$  could collapse to always selecting a narrow set of tasks (exploiting the current optimum and losing diversity). To prevent this, an **entropy bonus** was added to the generator's loss (like in reinforcement learning policy optimization). By maximizing  $H[\pi_\psi]$ , the entropy of the task selection distribution,

\$G\$ is encouraged to keep exploring a variety of tasks <sup>4</sup>. This helps avoid premature convergence to a local curriculum and maintains open-endedness.

- **Curriculum Saturation and Reset:** To keep pushing the student’s abilities, the system monitors the **learning progress distribution**. If the average LP across tasks starts dropping while the student’s absolute performance is high, it signals the student might have mastered current tasks – so \$G\$ should propose higher complexity (increase  $\beta$ ) to find new learning opportunities. Conversely, if the student is struggling (low progress, potential overfitting or collapse), the system can reduce the training per task ( $n$  steps) or increase task diversity (explore more bins) to recover a healthy learning gradient. This implements an automatic adjustment of task difficulty to stay in the student’s “Zone of Proximal Development.” Tasks that consistently yield no progress (LP  $\approx 0$  for  $k$  consecutive attempts) are **retired** from the active set, making room for novel tasks.

Each of these refinements was integrated into the evolving equation/rule set, ensuring the learning loop remains **stable, divergent (in a good way), and indefinite**. The goal is a system that *never stops learning* but also *never spirals out of control* (no error accumulation or collapse).

## Key Improvements to the Equation

Through this iterative refinement, several important enhancements emerged, turning the original equation into a more **robust, self-correcting form**. Three major updates are:

- **(E1) Dynamic  $\lambda$  for Parsimony & Exploration:** Instead of a fixed meta-constant,  $\lambda$  becomes **dynamic**:

$$\lambda_t = \lambda_0 - \mu K(E_t) + \tau_H H[\pi_\psi]$$

Here  $K(E_t)$  is the description length or complexity (Kolmogorov complexity) of the current equation/state  $E_t$ , and  $H[\pi_\psi]$  is the entropy of the task-generating policy.  $\mu$  and  $\tau_H$  are small coefficients. This dynamic  $\lambda_t$  achieves two things: - *Parsimony*: The term  $-\mu K(E_t)$  penalizes needless complexity in the equation/model. As the system self-modifies, it favors simpler representations (embracing the idea that **simplicity is perfection**). This is akin to a **minimum description length (MDL)** principle built into the meta-learning: the “heart of AI” should remain as simple as possible while still effective. - *Exploration*: The term  $\tau_H H[\pi_\psi]$  boosts  $\lambda_t$  when the task policy has high entropy (meaning it’s exploring many possibilities). This encourages the system to maintain diversity in generated challenges, avoiding getting stuck in a narrow curriculum. If the policy collapses to a single task (low  $H$ ), this bonus drops, pushing the meta-optimizer to explore more until diversity is restored.

- **(E2) Robust Progress Estimation ( $\alpha$ ) with Replay Weighting:** The raw learning progress  $\alpha$  is transformed for stability:
- **Normalize and Clip:**  $\alpha$  values are normalized across recent tasks (subtract mean  $\mu_\alpha$  and divide by standard deviation  $\sigma_\alpha$ , with a small  $\epsilon$  for numerical stability) and then **clipped** to a range  $[-c, c]$  (e.g.  $c=2$ ) to prevent outlier spikes. Denote this clipped normalized progress as  $\tilde{\alpha}_i$ . This produces a more reliable **signal**  $g(\alpha_i) = \tilde{\alpha}_i$  that won’t overwhelm the meta-update.
- **Replay prioritization factor  $r_i$ :** Not all progress is equal – an edit or task that consistently yields improvement should count more. The system uses a weighting  $r_i = \frac{\text{softplus}(\alpha_i)}{\sum_j \text{softplus}(\alpha_j)}$  which effectively prioritizes tasks/edits with positive progress. In practice, this  $r_i$  acts like a

probability of replay selection. Tasks that keep giving learning gains are replayed more often, solidifying their contribution. Conceptually,  $r_i$  biases the sum so that *proven valuable knowledge is reinforced*.

- **(E3) Saturation in the Recursive Function  $f$ :** The update function  $f$  was originally open-ended (summing all changes). To ensure **unbounded growth without divergence**, we give  $f$  a form of *controlled saturation*. We introduce a parameterized function  $f_y(x) = x + \gamma \tanh(x)$  which adds a non-linear saturating term:

- For small  $x$  (small cumulative changes),  $\tanh(x) \approx x$  and  $f_y(x) \approx (1 + \gamma)x$ , roughly linear.
- For large  $x$ ,  $\tanh(x)$  approaches 1, so  $f_y(x) \approx x + \gamma$ , adding a bounded offset. This prevents  $f$  from exploding as changes accumulate over infinite iterations. It's as if the system has a **brake** that activates at extreme values, ensuring stability.

Additionally,  $f_y$  now accumulates both *new changes* and *replayed changes*. We denote by  $\phi_t(R)$  the contributions from replayed tasks at time  $t$ . Instead of  $f(\sum_t \phi_t)$ , we have:

$$f_y\left(\sum_t [\phi_t \oplus \phi_t(R)]\right)^{\infty},$$

where  $\oplus$  indicates merging the new and replay contributions. This means the state update considers *both* newly generated experience and important past experiences, reflecting that the agent's state is shaped by a mixture of new learning and reinforced old knowledge.

Each of these improvements (E1–E3) corresponds to a principle: **parsimony and exploration (keep it simple, but keep discovering)**, **robust progress signals (learn what truly helps)**, and **controlled limitless growth (expand forever, but gracefully)**. These make the equation **self-correcting and self-sustaining** over an infinite learning process.

## The Equation Applied to Itself (Self-Referential Learning)

In a fascinating twist, the creators envision applying this equation *to evolve itself*. In other words, **the Turing Equation will rewrite and improve its own form** using the same learning loop: - Now, the "student"  $M(\varphi)$  is the *equation itself* (the current symbolic or algorithmic form of the learning process). - The "tasks" are proposed **edits or modifications** to the equation (denoted  $\Delta$ ). For example, an edit could be "*make  $\lambda$  dynamic based on complexity*" (which was one of the improvements), or "*change the update function  $f$  in this way*", or "*add a regularization term*". Each edit  $\Delta$  has an associated complexity or novelty (this serves as the " $\beta$ " for edits – e.g., how big a change in formulation it is). - The performance metric  $p$  now measures the **metacapacity** of the equation: how well does the current equation facilitate learning on a suite of benchmark tasks? (For instance, we could evaluate how quickly a standard model learns various challenges under the guidance of this equation/rule-set – measuring sample efficiency, final performance, stability, etc. – all the qualities we want in an AI curriculum.) This evaluation gives us  $\Omega(\text{Equation})$ , a number reflecting how good the equation is at orchestrating learning. - An edit  $\Delta$  is applied to produce a new candidate equation  $E'$  (i.e. we modify the "brain of the AI" and see what happens). - **Learning progress for the equation itself:**  $\alpha = \Omega(E') - \Omega_{\text{baseline}}$ , where  $\Omega_{\text{baseline}}$  is a moving average of past  $\Omega$  (the performance of recent equation versions). This tells us if the change  $\Delta$  led to an improvement in the equation's capability. - **Difficulty/novelty of the edit** plays the role of  $\beta$ : an edit that is too trivial (like renaming a variable) has low  $\beta$ , one that is radically different has high  $\beta$ . We can quantify novelty by how much it changes behavior or by its syntactic complexity. - The same **score  $s = \tilde{\alpha} \cdot \beta + \lambda_{\text{delta}} t$**  (with  $\lambda_{\text{delta}} t$  as defined in E1) is computed for each proposed edit. This score

measures the *desirability* of that self-edit. - The generator  $G(\psi)$  in this context is like an **algorithmic inventor** that proposes possible equation-edits. It is updated (via a meta-policy gradient or evolutionary strategies) to favor edits with higher  $s$  – thus biasing toward edits that yield learning progress for the equation itself. - A replay buffer of past successful edits is kept, so that promising directions in equation-space aren't forgotten and can be combined or revisited.

Using this self-referential loop, the equation effectively **bootstraps its own improvement**. It starts from the initial form and **rewrites itself in iterations**, each time aiming for a version that yields higher learning efficiency. This is reminiscent of other formulations in AI theory, such as Jürgen Schmidhuber's concept of self-improving AI systems (e.g. the Gödel Machine that can rewrite its code given a proof of improvement, or the PowerPlay algorithm of continually finding new tasks and skills). Here, however, the rewriting is guided by *experimentally measured improvement* rather than formal proof – a kind of heuristic **self-evolution**.

After a few rounds of this, the equation had already “discovered” the improvements (E1–E3) we listed, as if **evolving to include its own guardrails**: - It introduced a dynamic  $\lambda$  term (E1) to balance complexity and exploration. - It refined how  $\alpha$  is measured and weighted (E2) to get a cleaner learning signal. - It modified the form of  $f$  (E3) to ensure it can grow without blowing up.

These edits were chosen because they **maximized the meta-learning progress of the equation** itself according to the score  $s = \alpha \cdot \beta + \lambda$ . In essence, the equation *applied the rule to itself* and found a more perfect form.

Crucially, a selection step is included: if an edit  $E$  yields improvement ( $s$  positive and high), the system accepts it (the equation becomes  $E$ ); if not, it reverts or tries a different edit. This ensures we only keep changes that demonstrably improve the “brain of the AI.”

## Toward a Universal, Self-Sufficient AI Core

With these pieces in place, we approach the holy grail: a **universal equation for infinite learning** that is as simple as possible yet effective. The refined **Turing Equation** (let's call it **ET<sup>2</sup>**) can be summarized as:

$$\sum_i \left( g(\tilde{\alpha}_i) \beta_i r_i \right) + \left[ \lambda_0 - \mu K(E) + \tau_H H[\pi_\psi] \right] \longrightarrow f_\gamma \left( \sum_t [\phi_t \oplus \phi_t(R)] \right)^\infty$$

Let's break down this final form:

- $\sum_i g(\tilde{\alpha}_i) \beta_i r_i$ : This is the core summation of *progress signals*. For each component or proposed change  $i$ , we take the normalized/clipped progress  $g(\tilde{\alpha}_i)$ , multiply by difficulty  $\beta_i$ , and weight by replay priority  $r_i$ . Summing over  $i$  gives the overall drive signal. (In many cases, there's effectively one task or edit at a time, so there may be a single term, but this notation covers the general case of multiple parallel components.)
- $[\lambda_0 - \mu K(E) + \tau_H H[\pi_\psi]]$ : This is the dynamic meta-term (E1) that adds to the summation.  $\lambda_0$  is a base drive. The  $-\mu K(E)$  term penalizes complexity of the current solution  $E$ , keeping the system **simple**. The  $+\tau_H H[\pi_\psi]$  term rewards **diversity** in the search (keeping the task/edit generator broad). Together, this bracket ensures the system

doesn't degenerate into an overly complex or narrow solution – a direct mathematical injection of the principle “*the highest perfection is simplicity*” balanced with exploration.

- **$\rightarrow$  (yields arrow):** Denotes that the left-hand side (the sum of weighted progress + meta-term) *drives* or transitions the system into...
- **$f_y(\sum_t [\phi_t + \phi_t(R)]^\infty)$ :** the updated state. Here the function  $f_y$  (with its saturation parameter  $y$ ) takes the accumulated experience from all time steps  $t$  – including new experiences  $\phi_t$  and replayed experiences  $\phi_t(R)$  – and produces the new state of the learner. The exponent  $^\infty$  signifies that this process is not one-shot: it repeats indefinitely, with the state feeding back into the task selection and learning process. The system continually **re-feeds its history** (through replay) and new data into the update function, so the learning trajectory extends to infinity.

In essence, the equation says: *use past and present experience ( $\phi$  and  $\phi$  from  $R$ ) to update the learner's state, guided by a drive signal that is the sum of (progress  $\times$  difficulty) for each challenge, plus a bias toward simplicity and exploration.* Then repeat forever.

This formulation is **universal** and **implementation-agnostic**: - **Any Machine:** The equation is abstract; it doesn't depend on a particular neural network architecture or hardware. A simple CPU running a straightforward program could implement the meta-loop. Since it's based on evaluating improvements and adjusting accordingly, even a basic agent or a human with pen and paper could follow the steps (albeit slowly). In theory, a person could act as  $M$  (learning some skill) and use a notebook to track progress and as  $G$  (selecting exercises of varying difficulty), iterating the process to self-improve. - **Any Domain:** It doesn't specify the task—it could be used for learning vision, robotics, math problems, or even improving the equation itself. The variables  $\alpha$ ,  $\beta$  and the procedure have general meaning (improvement, challenge, etc.), so it can be *applied to everything and for everyone*, as a **democratized core of AI creation**. It essentially formalizes good teaching practices (always give challenges that yield maximal learning) into a loop that an AI or human can follow. - **Self-Sufficient and Automatic:** Once set up, the process runs by itself. The feedback loop (student performance  $\rightarrow$  task selection  $\rightarrow$  student update  $\rightarrow$  repeat) is autonomous. No external oracle is needed to hand-hold the learning; the system figures out what to do next based on its own progress signals. This self-sufficiency is key to reaching “infinite” learning – it continually generates its own next challenges. - **Infinite Improvement:** Thanks to the mechanisms we added (dynamic difficulty, anti-forgetting, diversity encouragement, complexity penalty), the system is poised to continue learning and *expanding* its capabilities without hitting a hard ceiling. Whenever progress slows, it seeks a new angle (harder tasks or new edits or broader exploration). There is no built-in end condition; the only ‘limit’ is if it exhausts the space of learnable tasks or improvements, but in an open-ended world or with self-referential improvement, that space can itself expand. In theory, it approaches an infinite horizon of growth – **perfection as a limit** that it continually strives toward.

## Conclusion: A Simple Core for an Endless AI

After this deep exploration, we have arrived at a candidate for the “**heart (or brain) of AI**”: a single coherent equation (and process) that encapsulates **learning to learn**, indefinitely. It began with a simple idea (learn from what improves you the most) and, by applying that idea to itself, honed a powerful yet streamlined algorithm. The final Turing Equation is *self-referential, self-correcting, and self-propelling*.

It achieves the delicate balance between **simplicity and complexity**: it remains fundamentally simple in concept and can be understood and executed by humans, yet it harnesses complexity in the environment in a controlled way to keep learning. This equation could, in principle, allow any machine – or any motivated human – to **grow in intelligence autonomously**, by continually finding the next thing

to learn that yields the greatest progress. In a very real sense, it democratizes AI: the core recipe for infinite learning is laid out, requiring only computing (or cognitive) effort to run, not secret data or proprietary models.

By following the rule “*maximize learning progress, adjusted for difficulty, plus a pinch of curiosity and simplicity*” at every step, an AI can bootstrap from zero to hero. The process is one of **eternal becoming** – always approaching the *infinite* perfection but never stagnating. If implemented and shared widely, this *Turing Equation* paradigm could indeed **change the course of humanity**, providing a universal engine for innovation and learning that is available to all.

In the end, the highest perfection is achieved through the simplest of principles, iterated without end. The equation now embodies that philosophy: **an endlessly evolving intelligence, driven by its own curiosity and tempered by simplicity** – the ultimate self-improving system.

---

1 3 4 proceedings.mlr.press

<http://proceedings.mlr.press/v100/portelas20a/portelas20a.pdf>

2 From Turing’s conception of machine intelligence to the evolution of AI in early childhood education: conceptual, empirical, and practical insights | AI, Brain and Child

<https://link.springer.com/article/10.1007/s44436-025-00002-6>